

Random Forest Model for Primer Presence Prediction

Purpose

The Random Forest (RF) model is designed to predict the presence of a primer in a given SRA sequencing dataset by analyzing the quality score distribution of the initial subset of reads.

Data Source

- A curated collection of SRA runs known to contain or lack specific primer sequences.
 - Each SRA run was pre-processed with `fastp` to trim adapters and filter low-quality reads.
 - From each processed sample, the **first 1000 reads** were extracted using `fastq-dump` with `-X 1000`.
-

Feature Engineering

Two sets of quality scores were extracted from each sample:

- **s1**: The quality scores from positions 1–5 of the first 1000 reads.
- **s2**: The quality scores from positions 6–10 of the same reads.

From each of these two segments, the following 8 statistical features were computed:

1. `count` – Number of quality scores extracted
2. `mean` – Average quality score
3. `median` – Median quality score
4. `std` – Standard deviation
5. `min` – Minimum value
6. `max` – Maximum value
7. `skew` – Skewness of the distribution (25th percentile)
8. `kurtosis` – Kurtosis of the distribution (75th percentile)

Total of **16 features** per sample:

- `1_5_count, 1_5_mean, ..., 1_5_kurtosis`
 - `6_10_count, 6_10_mean, ..., 6_10_kurtosis`
-

Labels

Each sample was labeled as:

- **1** if the primer was confidently present
 - **0** if the primer was confidently absent
-

Model Training

- **Algorithm:** RandomForestClassifier from `scikit-learn`
 - **Version:** `scikit-learn` 1.2.1
 - **Parameters:**
 - `n_estimators=100`
 - `random_state=42`
 - **Training/Test Split:** 80/20 split with stratification
 - **Scaler:** None (RF handles raw value ranges robustly)
-

Model Persistence

- The trained model was serialized using `joblib.dump(model, "rf_model.pkl")`
 - Later used via `joblib.load("rf_model.pkl")` within the `HVRegLocator` script
-

Usage in Pipeline

- When the `--model` flag is passed, quality scores are extracted from trimmed FASTQ reads
 - Features are computed and passed to the loaded model
 - Prediction is stored as:
 - `Primer_Presence: "Yes" / "No"`
 - `Score_Primer_Presence: Probability score from predict_proba()`
-

Version Compatibility Note

The model was trained with **scikit-learn 1.2.1**, and should ideally be used in the same version to avoid `InconsistentVersionWarning` during unpickling.

Model Accuracy: 99.96%

Precision:

- No primers: 99.96%
- With primers: 100%

Recall:

- No primers: 100%
- With primers: 99.55%